

eBPF 101

“ eBPF does to Linux what JavaScript does to HTML ”
Brendan Gregg

Agenda

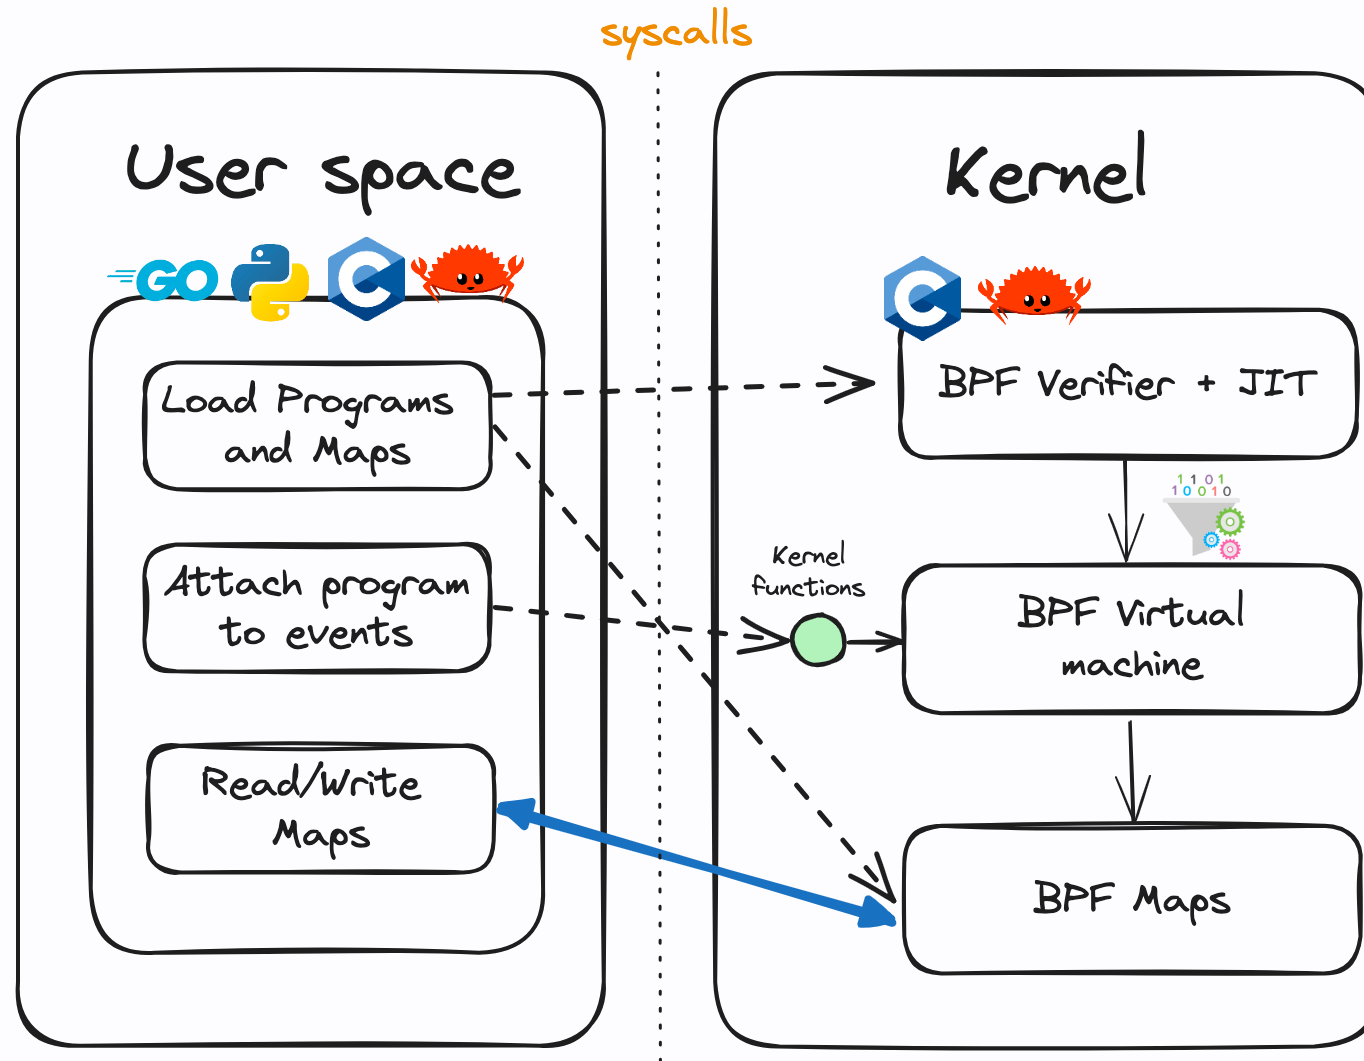
- What is eBPF?
- Why do we need eBPF?
- Examples
 - The EASY way
 - The HARD way
- Takeaways



What is eBPF?

- BPF, **Berkeley Packet Filter**, was introduced to Linux in 1997, kernel version 2.1.75
- It has been used in the tcpdump utility to capture packets
- In 2012 `seccomp-bpf` was added to kernel version 3.5, allowed to take decisions about user apps making system calls
- eBPF, **extended BPF**, started in 2014 with new features...

The eBPF basics



Why do we need eBPF?

- To add new functionalities to the Kernel without the overhead of creating Kernel Modules, and with **dynamic loading**
- **Better performance** skipping the cost of transitioning between kernel and user space (for each event)
- It's a great match for **cloud native** environments, no more need for sidecar pattern

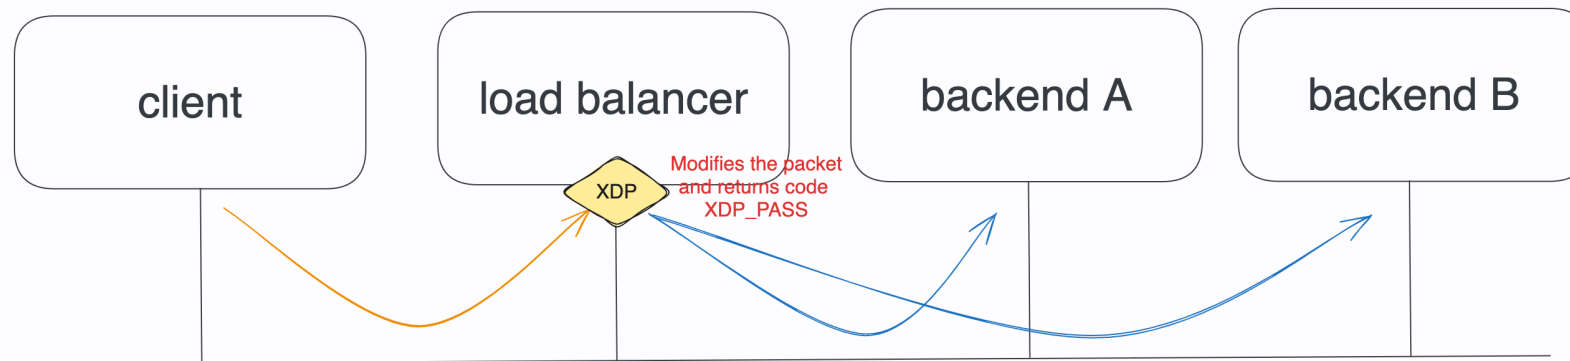
Where to use eBPF?

Observability, Networking & Security

- App observability (Beyla, Tetragon)
- Profile & Tracing (bpftrace, sysdig)
- Load-balancing (Katran)
- Network monitoring (Netobserv, pwru)
- Kubernetes (Cilium, Calico)
- Firewalls (Bpfilter)
- DDoS mitigation (Cloudflare, Meta)

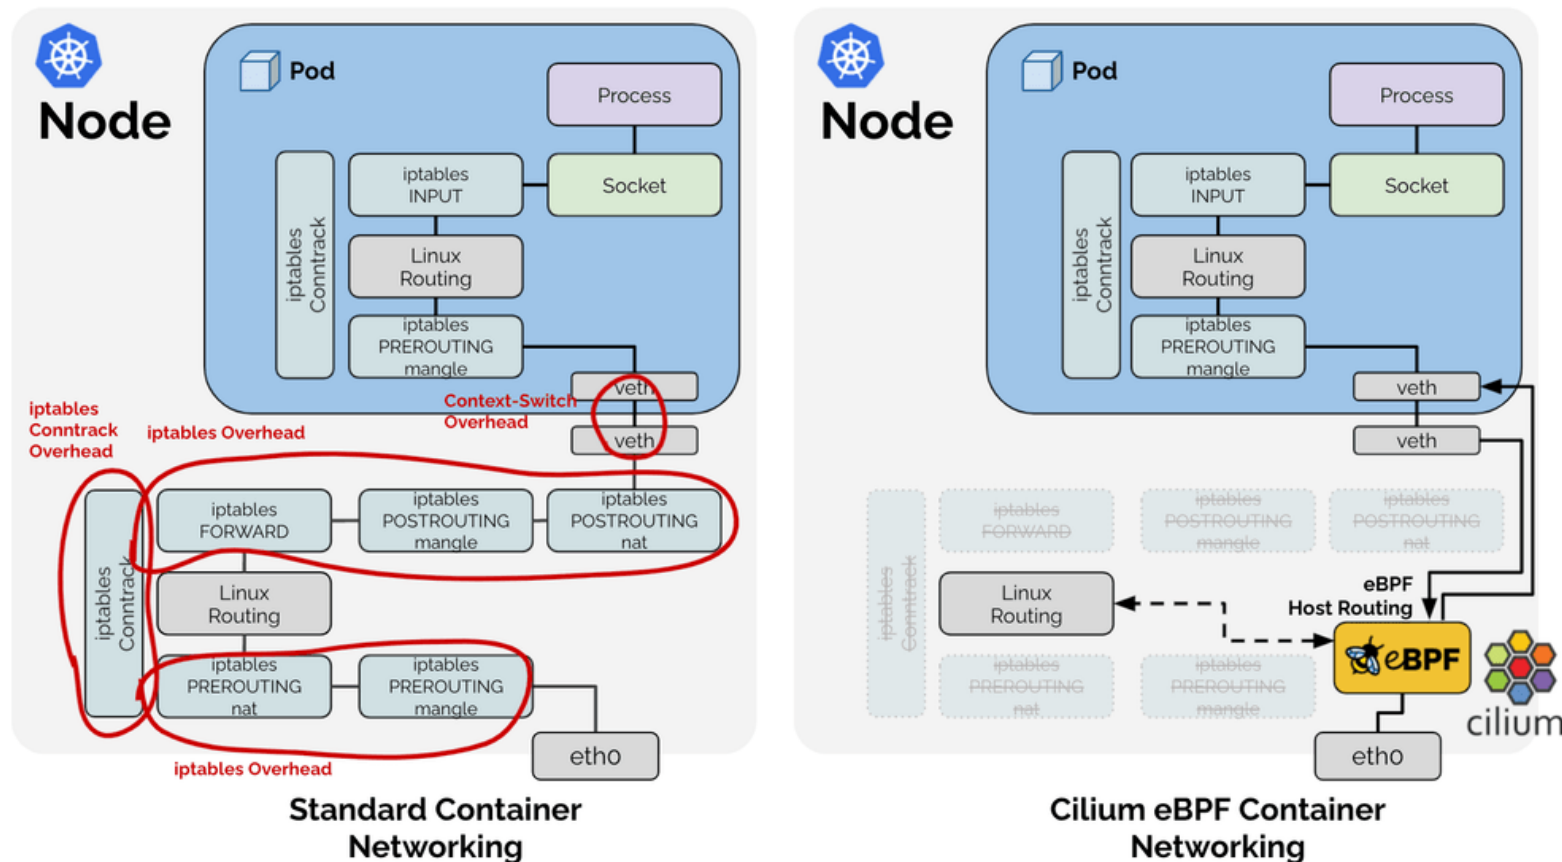
Tools available: <https://ebpf.io/applications/>

A Load Balancer from scratch



https://www.youtube.com/watch?v=L3_AOFSNKK8

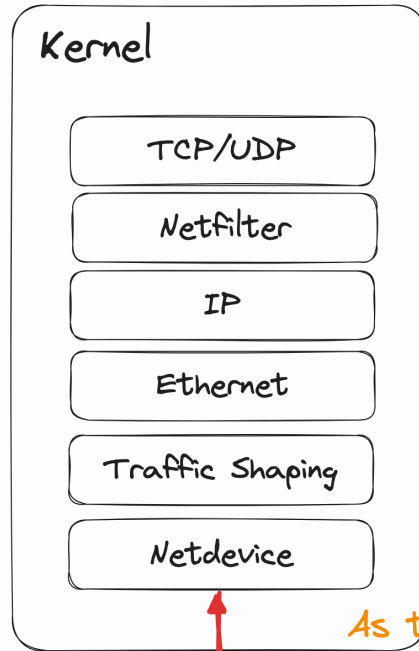

eBPF in Kubernetes



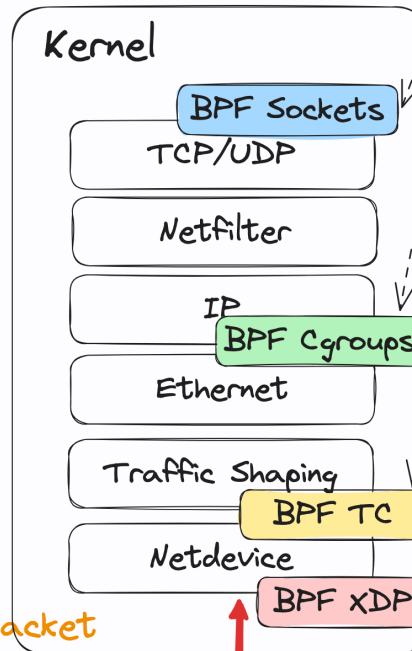

<https://cilium.io/blog/2021/05/11/cni-benchmark/>

Comparing Linux, VPP, and eBPF



Linux



eBPF



DPDK & VPP



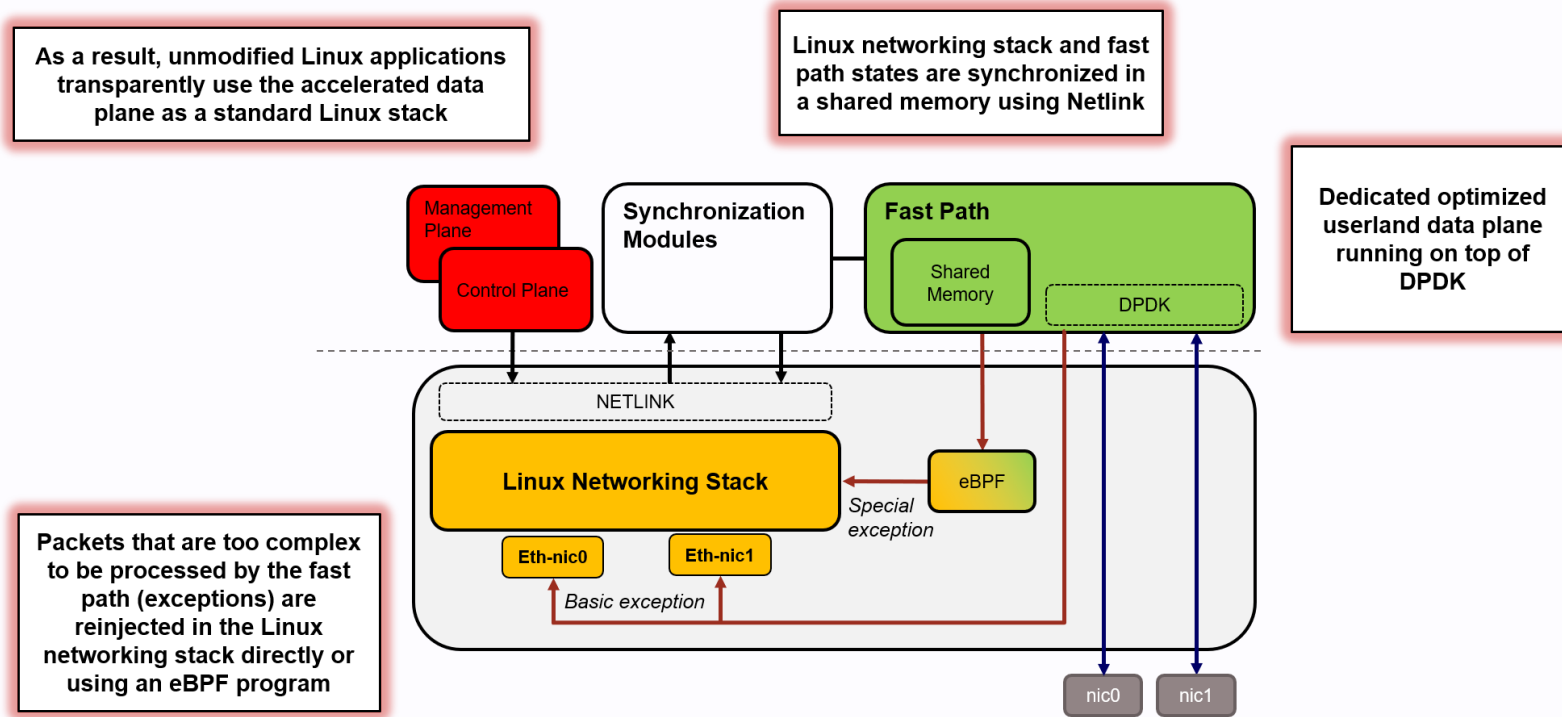
As the packet traverses the stack, you get more functions, but less pps

NIC

Some numbers

| Linux | DPDK & VPP | XDP and eBPF |
|--|---|--|
| ~ 30 cores for 28Mpps forwarding | ~ 3 cores for 28Mpps forwarding | ~ 5 cores for 28Mpps forwarding |
| Rich feature set. Socket programming as usual (IPtables) | Amazing performance!, but you lose the NIC from the Kernel and need to reimplement everything | Excellent performance with early packet processing and fallback to Kernel. Access to FIB table (<code>bpf_fib_lookup</code>) |
| Optimized for general purpose | Great for forwarding plane, with networking features | Optimized for best of both worlds |

Both can cooperate



<https://www.6wind.com/6windgate-5-0-versus-vpp-fast-path-benefits-from-linux-integration-and-ebpf/>

Examples

Check the main [README.md](#) for details to setup the environment and run the examples

The EASY way

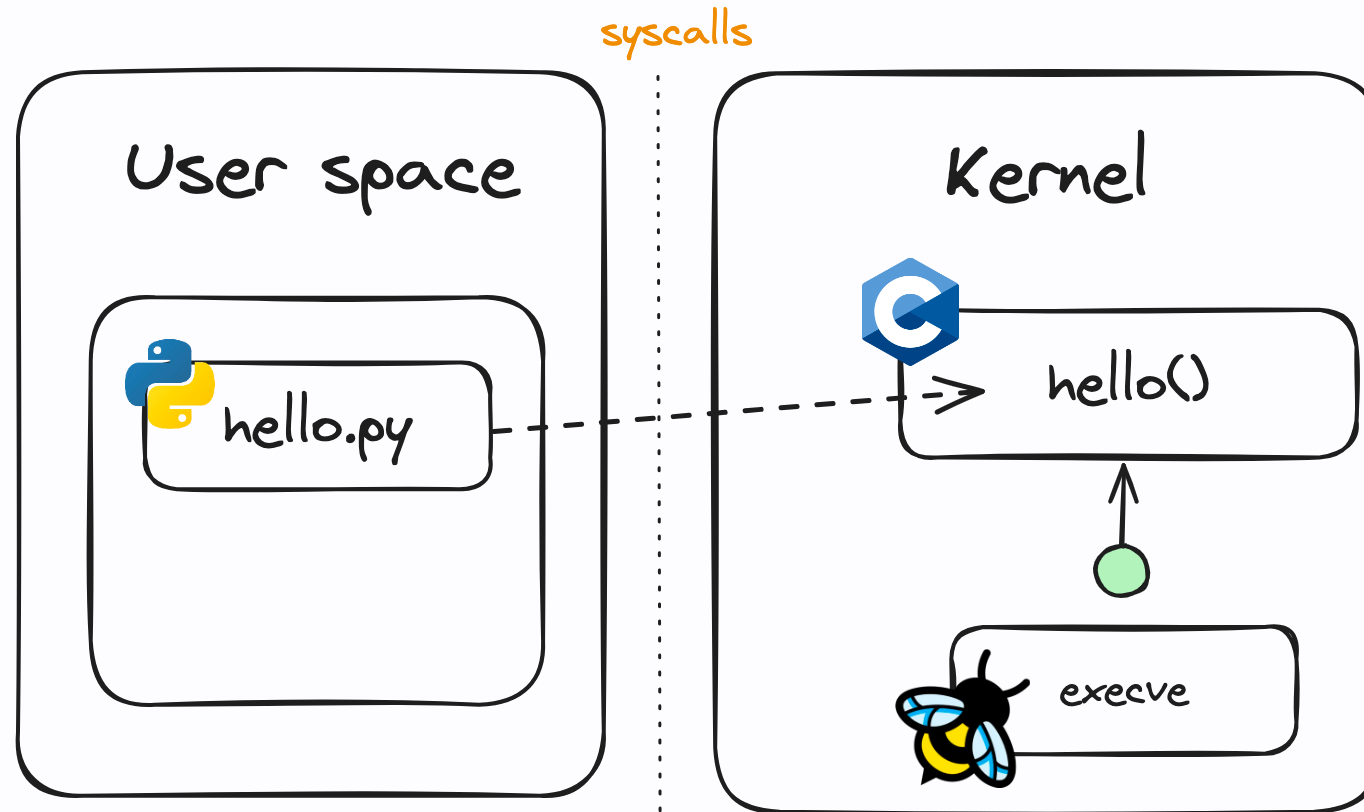
The BCC library, BPF Compiler Collection (BCC)

“ BCC is a toolkit for creating efficient kernel tracing and manipulation programs with eBPF, and includes several useful tools and examples. ”

- kernel instrumentation in C (and includes a C wrapper around LLVM)
- front-ends in Python and lua
- takes care of a lot of low-level details

There are frameworks for other languages (`ebpf-go`)

Hello World



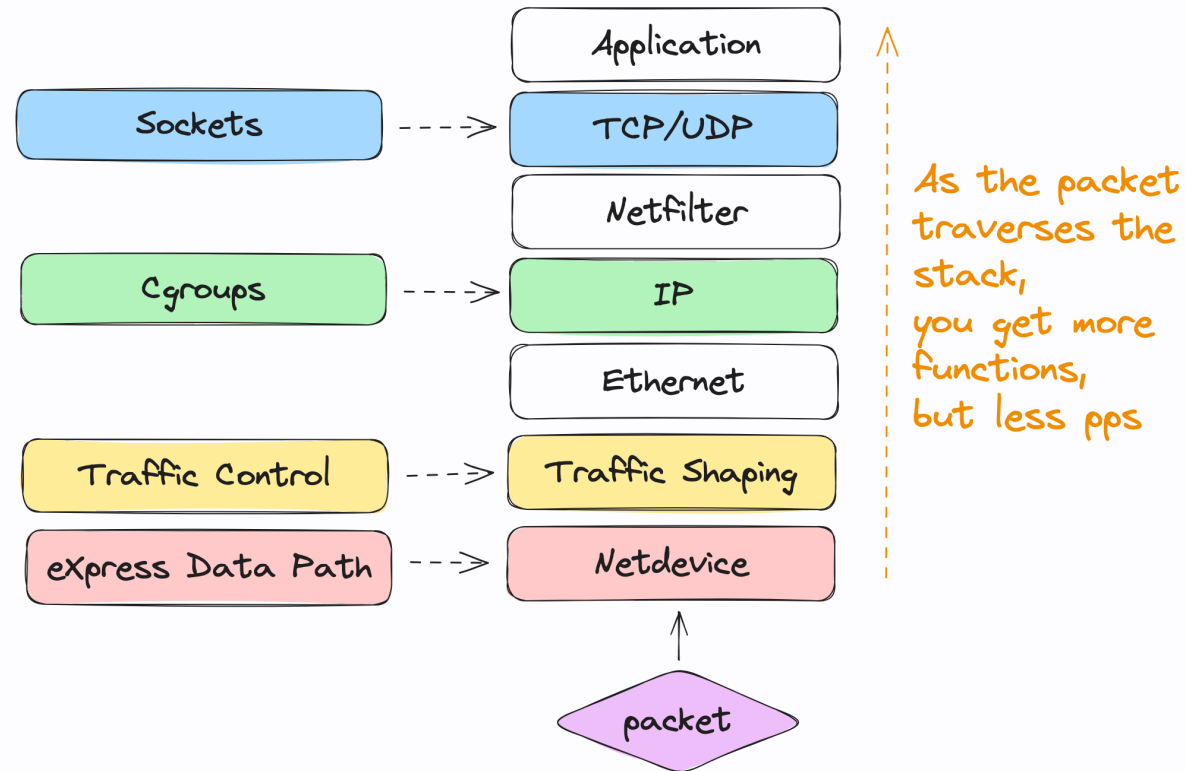
Code available at <https://github.com/chadell/ebpf-playground/blob/main/examples/python/hello.py>

Hello World - code

```
from bcc import BPF
program = r"""
int hello(void *ctx) {
    bpf_trace_printk("Hello NetBCN!");
    return 0;
}
"""

# Loads the C code from a string
b = BPF(text=program)
# Retrieves the syscall for execve
syscall = b.get_syscall_fnname("execve")
# Attaches the C function hello to the event
b.attach_kprobe(event=syscall, fn_name="hello")
# Outputs the trace in the screen
b.trace_print()
```

The ping example



Code available at <https://github.com/chadell/ebpf-playground/blob/main/examples/python/ping.py>

The ping example code

```
from bcc import BPF
program = r"""
#include "network.h"
#include <bcc/proto.h>
#include <linux/pkt_cls.h>
int xdp(struct xdp_md *ctx) {
    void *data = (void *)(long)ctx->data;
    void *data_end = (void *)(long)ctx->data_end;
    if (is_icmp_ping_request(data, data_end)) {
        bpf_trace_printk("Got an ICMP packet");
        return XDP_PASS;
    }
    return XDP_PASS;
}
"""

b = BPF(text=program)
fx = b.load_func("xdp", BPF.XDP)
BPF.attach_xdp("lo", fx, 0)
b.trace_print()
```

The HARD way

Let's use C code without helpers:

- Use the `bpftool`
- Compile the C program
- Load the eBPF code
- Attach to the loopback XDP hook

Code available at <https://github.com/chadell/ebpf-playground/blob/main/examples/c/hello.bpf.c>

C BPF code

```
#include <linux/bpf.h>
#include <bpf/bpf_helpers.h>
// This global variable it's converted to a MAP
int counter = 0;
// SEC() is a macro that defines the type of eBPF program
SEC("xdp")
int hello(struct xdp_md *ctx) {
    // bpf_printk is a helper function by libbpf
    bpf_printk("Packet received %d", counter);
    counter++;
    // XDP functions return the action for the packet
    return XDP_PASS;
}
// The eBPF verifier inspects the license of eBPF programs
char LICENSE[] SEC("license") = "Dual BSD/GPL";
```

C BPF step by step

- Install `libbpf`
- Compile the program (`-target bpf`)
- Check the compiled object
- Install `bpftool` to
 - Load the eBPF program and check it
 - Attach the eBPF program to an interface
 - Inspect the eBPF Maps
- See what's going on in the `trace_pipe`
- Detach and unload the eBPF program, with `bpftool`
- Let's make the Verifier unhappy

Takeaways

- The Verifier can get complicated
- Some kernel familiarity is needed
- Even there are some helpers, there are some limitations
- Kernel support has to be assessed

Resources

- Book: [Learning eBPF](#), by Liz Rice
- <https://ebpf.io/labs/>
- <https://gist.github.com/satrobit/17eb0ddd4e122425d96f60f45def9627>
- <https://speakerdeck.com/fedepaol/ebpf-for-the-rest-of-us-golab-2023>

Thanks!

Contact via LinkedIn [@christianadell](#)